CC LAB 2 - MONOLITHIC ARCHITECTURE

Name: Smruthi B S

PES1UG23AM308

SEC- F

19-01-2026

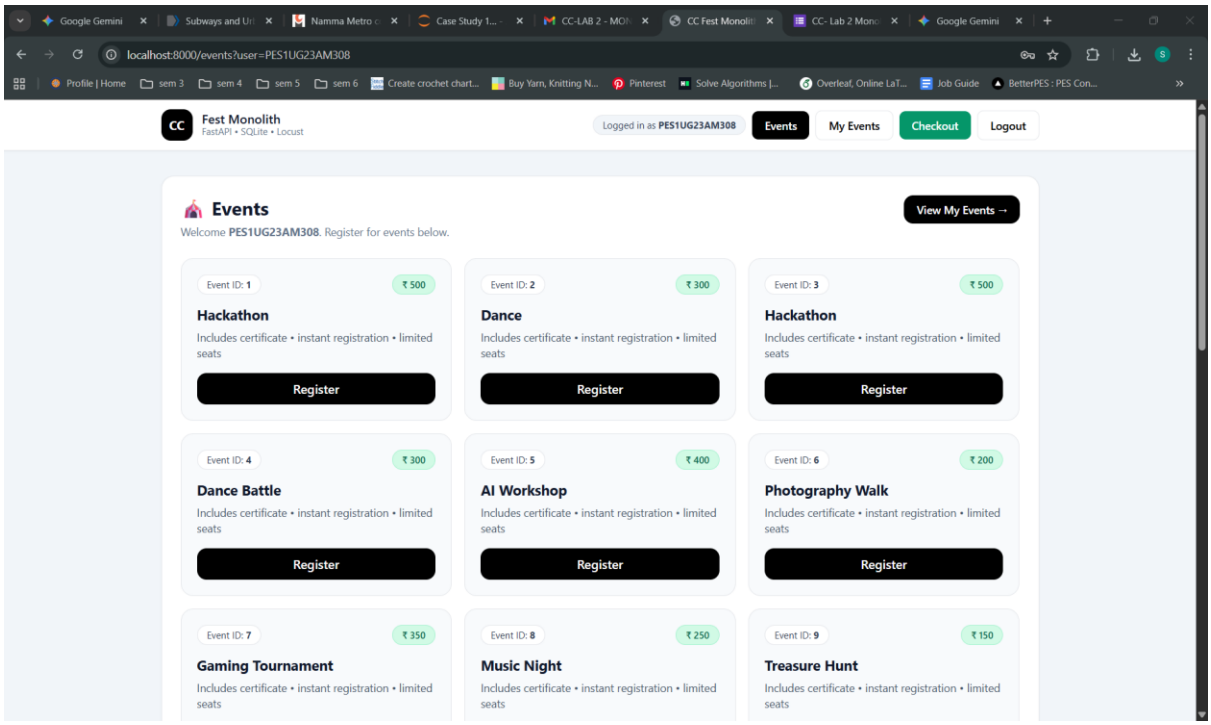

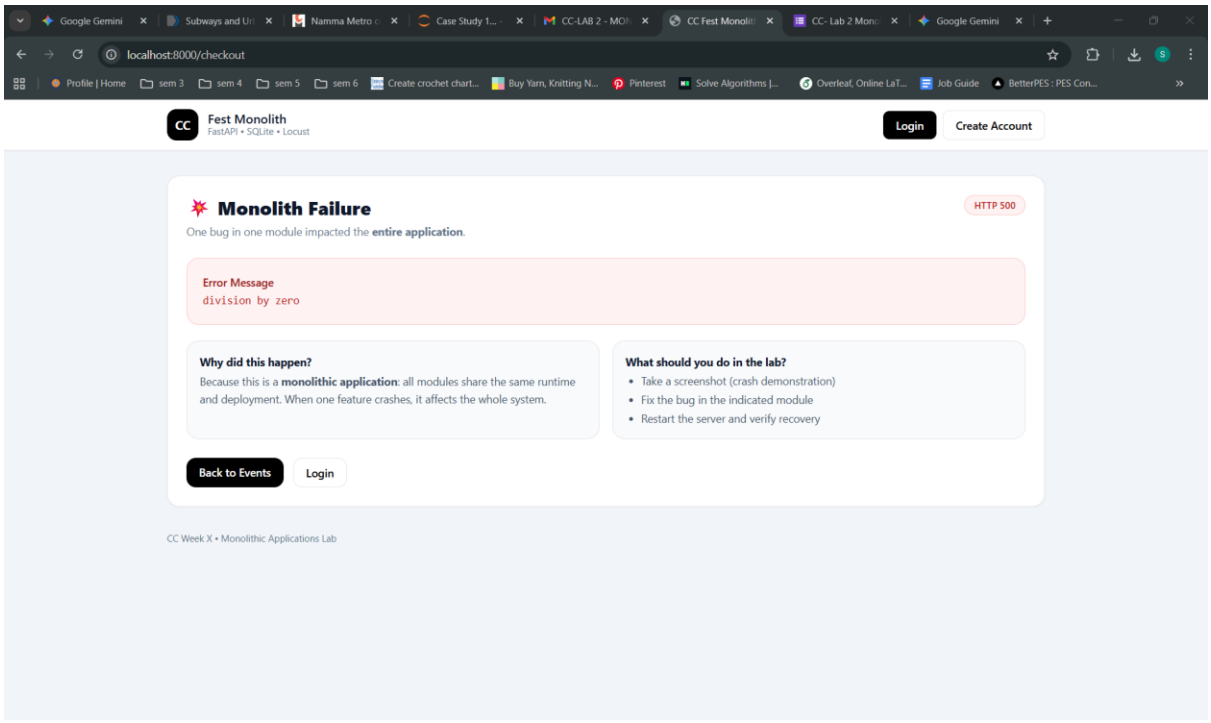
localhost:8000/checkout
Fest Monolith
FastAPI • SQLite • Locust
Login
Create Account
Monolith Failure
HTTP 500
One bug in one module impacted the entire application.
Error Message
division by zero
Why did this happen?
Because this is a monolithic application: all modules share the same runtime and deployment. When one feature crashes, it affects the whole system.
What should you do in the lab?
Take a screenshot (crash demonstration)
Fix the bug in the indicated module
Restart the server and verify recovery
Back to Events
Login
CC Week X • Monolithic Applications Lab

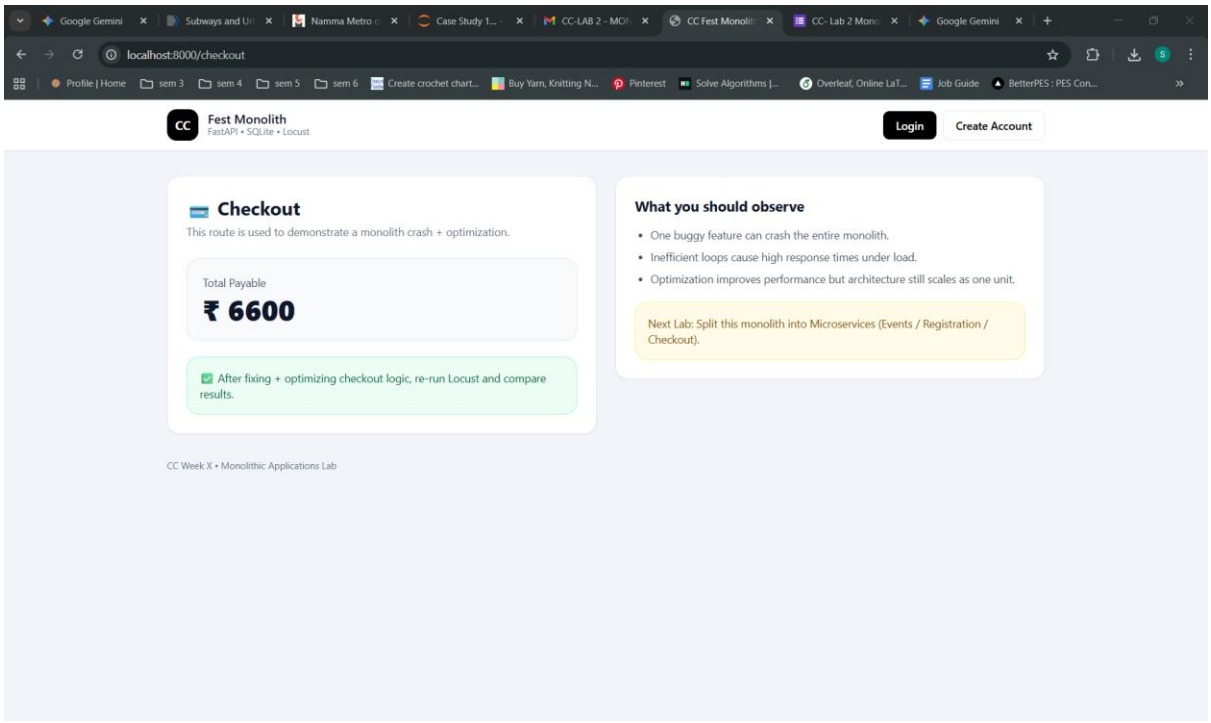
localhost:8000/checkout
Fest Monolith
FastAPI • SQLite • Locust
Login
Create Account
Checkout
This route is used to demonstrate a monolith crash + optimization.
Total Payable
₹ 6600
After fixing + optimizing checkout logic, re-run Locust and compare results.
What you should observe
One buggy feature can crash the entire monolith.
Inefficient loops cause high response times under load.
Optimization improves performance but architecture still scales as one unit.
Next Lab: Split this monolith into Microservices (Events / Registration / Checkout).
CC Week X • Monolithic Applications Lab

PART 6: Optimize the Checkout Route

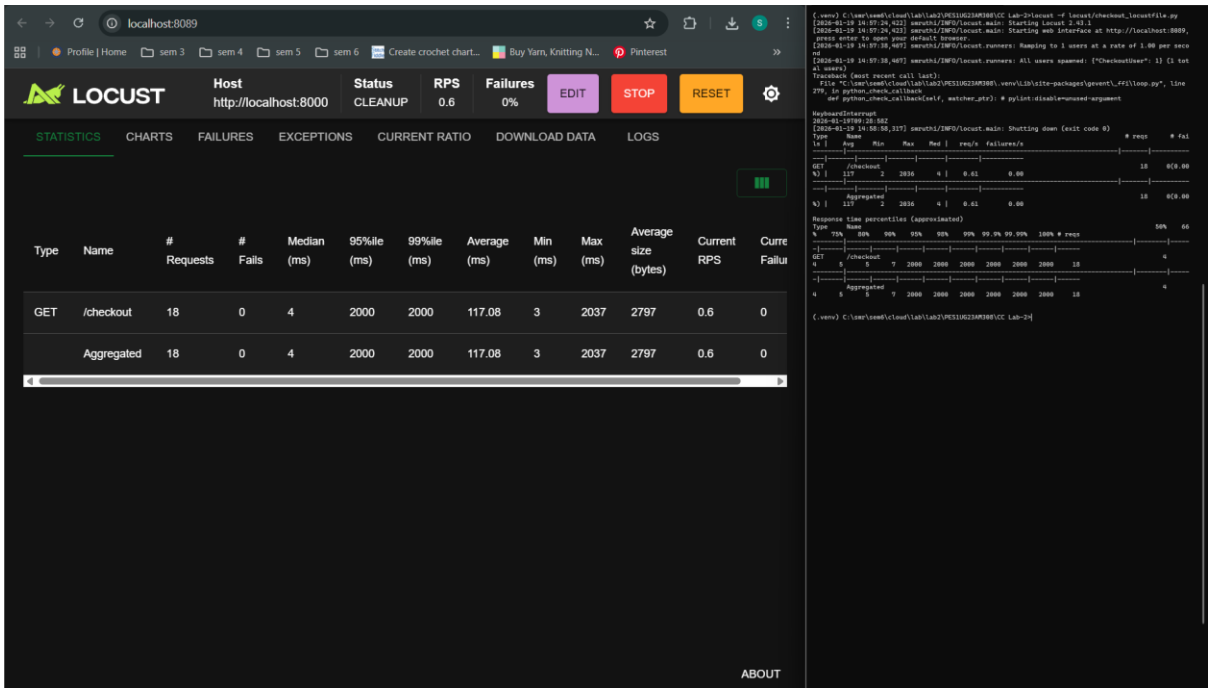

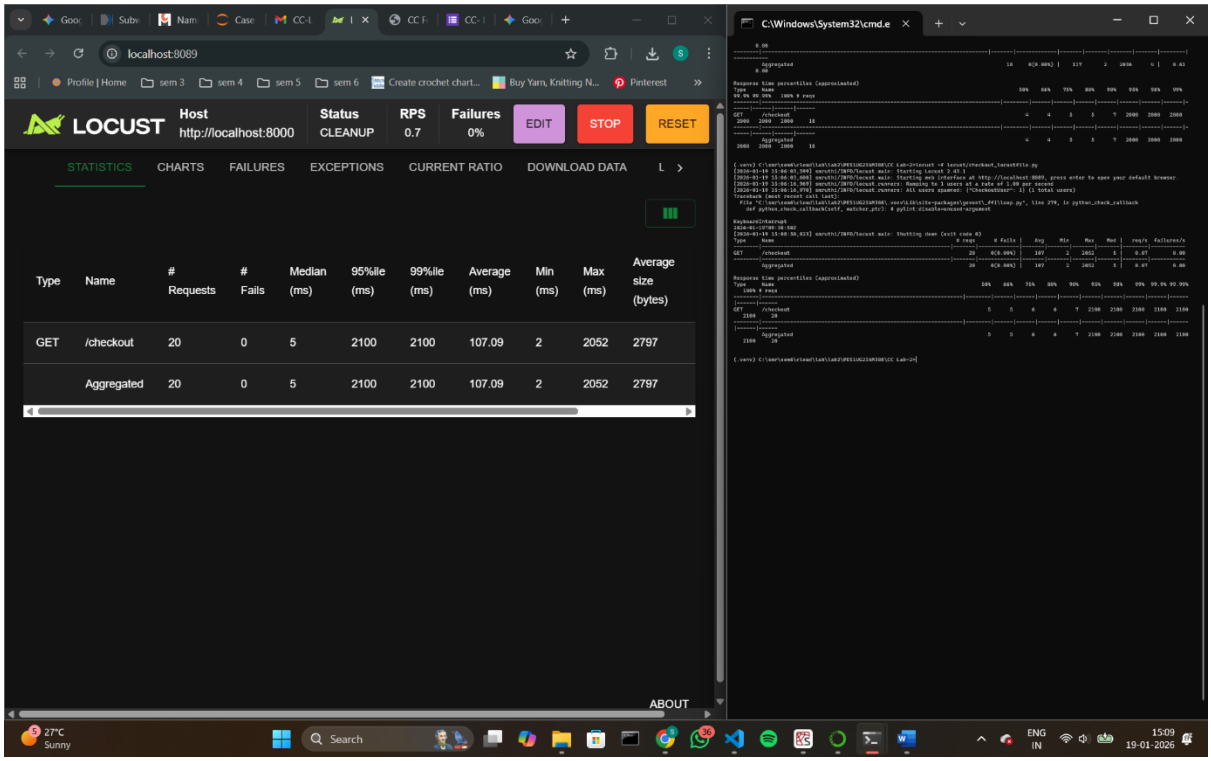

Route 1: /events

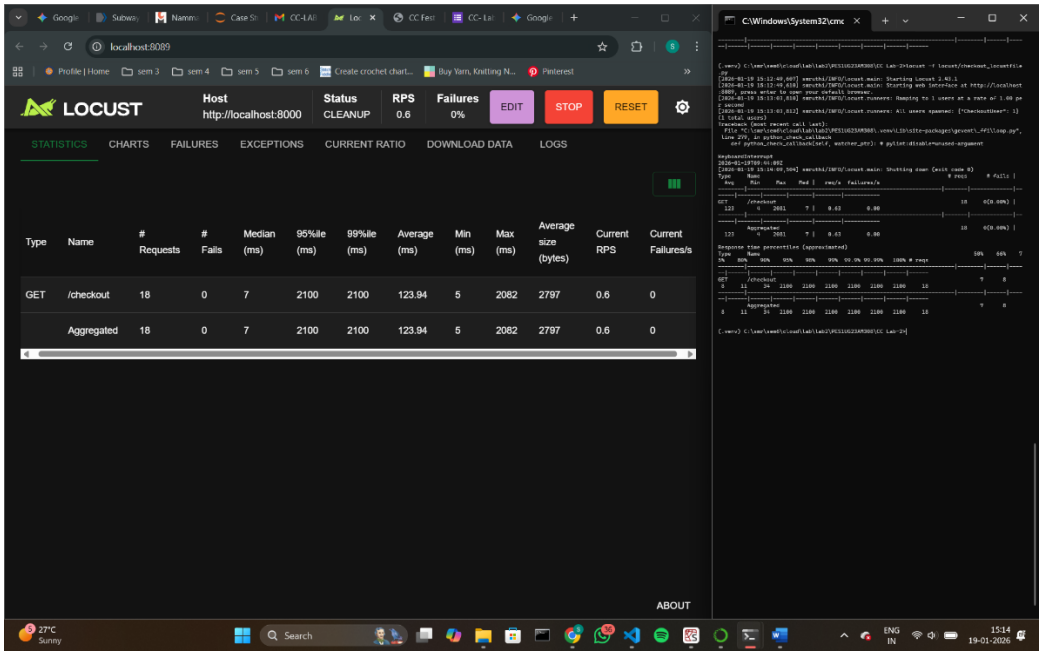

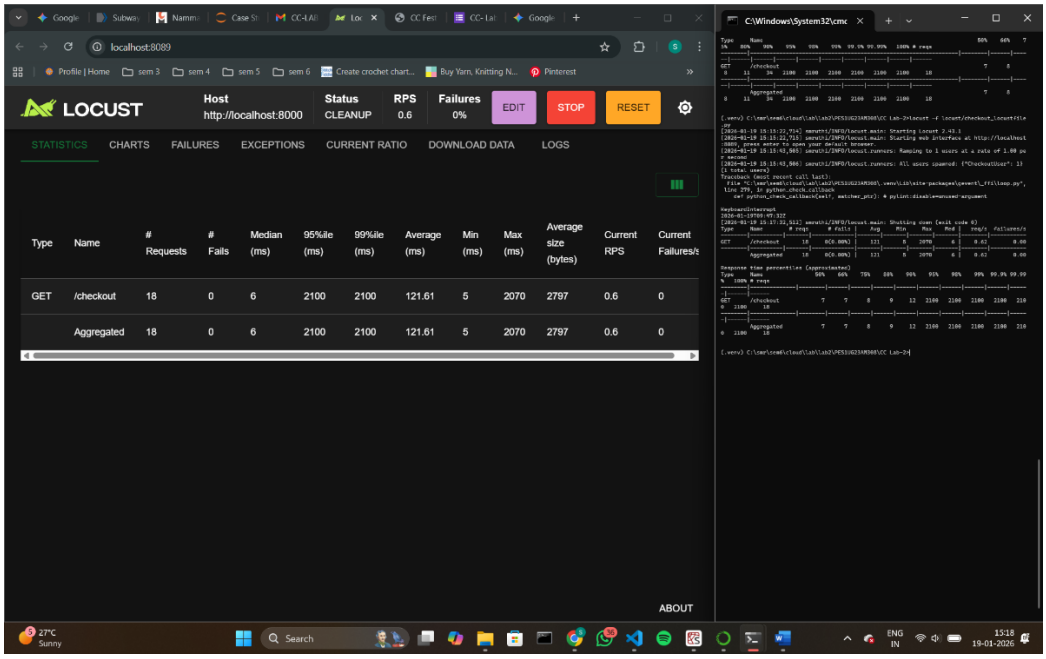

The bottleneck: The bottleneck was a computational overhead caused by a large for loop performing millions of modulo operations. This blocked the event loop, preventing the server from handling concurrent requests efficiently.

C**hanges made:** I ensured the "waste" loop (for i in range(3000000): waste += i % 3) was removed or remained commented out. I also verified that the database connection is handled efficiently without unnecessary re-querying.

Why the performance improved: Removing the loop shifted the operation from CPU-bound (slow) to I/O-bound (fast). The server no longer spends millions of cycles on useless math, allowing it to respond immediately after fetching data from the database.

Route 2: /my-events

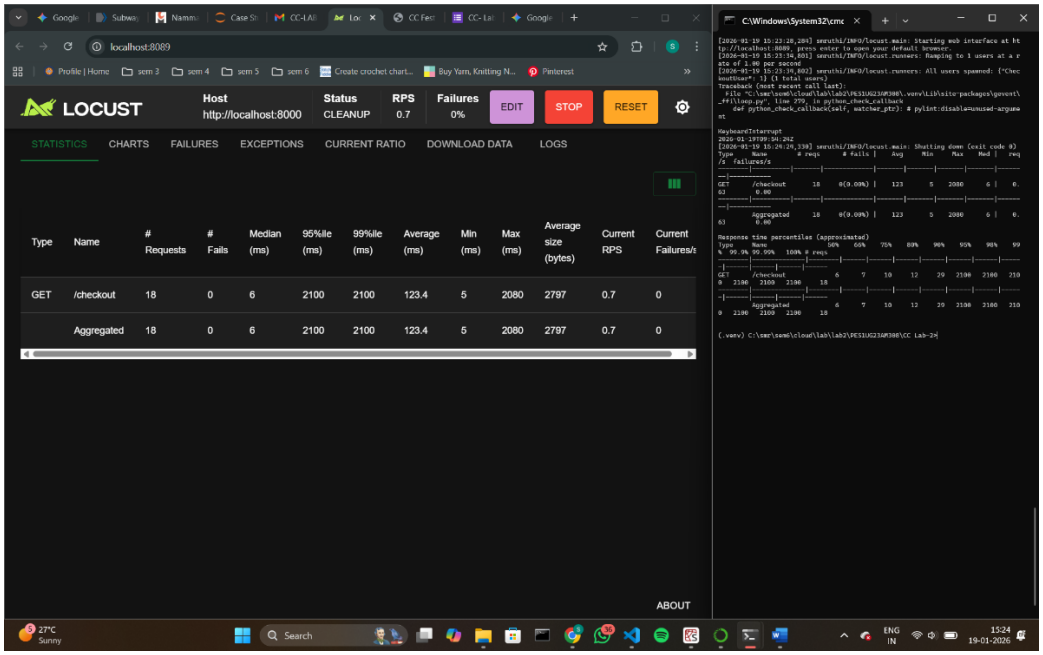

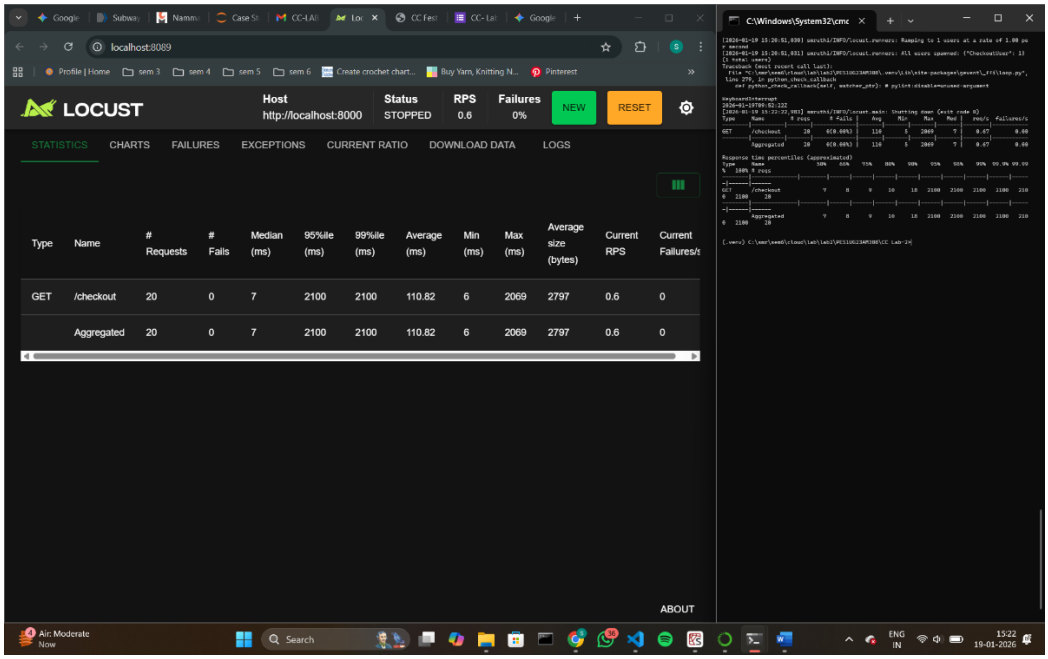

The bottleneck: The bottleneck was an artificial delay created by a dummy counter loop (for _ in range(1500000): dummy += 1). This caused high latency for every user trying to view their registered events.

Changes made: I deleted the dummy loop and the associated logic. I also ensured the SQL JOIN query is as lean as possible to minimize the time the database spends searching the registrations table.

Why the performance improved: The performance improved because the execution time per request dropped significantly. By eliminating the manual loop, the response time is now dictated only by the speed of the SQL query, leading to higher Requests Per Second (RPS) in Locust.